

Homework 2

Implement the Canny edge detector for a grayscale image.

Step 1. Write a function, called `my_Normalize(img)`, to prepare the image:

- (1) If `img` is a grayscale image, do nothing; if it is a color image, convert it to a grayscale image;
- (2) Normalize this image so that it becomes a matrix whose elements are float-point numbers and within range $[0,1]$

Step 2. Write a function, called `my_DerivativesOfGaussian(img, sigma)`, to compute the derivatives of smoothed images:

- (3) Construct two 3×3 Sobel kernels, S_x and S_y as the derivative operators on both x and y directions
- (4) Construct a Gaussian kernel g_σ where σ is the variance specified as the input parameter `sigma`; set its mask size to $6\sigma + 1$ (Note: Check example program 9)
- (5) Construct the derivative of Gaussian kernels, g_x and g_y by convolving the above two kernels: $g_x = S_x * g_\sigma$; $g_y = S_y * g_\sigma$
- (6) Apply both kernels g_x and g_y on the input image, and get two resultant images I_x and I_y
- (7) Normalize both I_x and I_y and render them
- (8) Return the **un-normalized** derivative images I_x and I_y

Step 3. Write a function, called `my_MagAndOrientation(Ix, Iy, t_low)`, to compute the magnitude and orientation of the gradient image:

- (9) Compute the magnitude image: $m(x, y) = \sqrt{I_x(x, y)^2 + I_y(x, y)^2}$;
- (10) Compute the orientation of gradient, store it in an image: $O(x, y) = \tan^{-1}(I_y(x, y)/I_x(x, y))$ (Note: you can use `numpy.arctan2`. See <https://docs.scipy.org/doc/numpy/reference/generated/numpy.arctan2.html>)
- (11) Round each element in O into one of the four values: $[0, \frac{\pi}{4}, \frac{\pi}{2}, \frac{3\pi}{4}]$ (check Lecture 10 slides page 7). Here save it as one of the four integers: $[0, 1, 2, 3]$, respectively.

Note: In the slides, the y^+ axis is upwards, while in a matrix representation, the y^+ direction is downwards. An easy modification to tackle this is to switch case 1 and case 3 in your implementation. Check your computation using 'TestImg1.jpg' example.

- (12) Normalize the magnitude image.
- (13) Plot the normalized magnitude image; and the rounded integer orientation image; return them.

Step 4. Write a function, called `my_NMS(mag, orient, t_low)`, to perform non-maximal suppression along the gradient direction:

- (14) Create a zero matrix called `mag_thin`
- (15) For each pixel `mag[i][j]` whose value is smaller than `t_low`, ignore it
- (16) For each pixel whose value is bigger than `t_low`, if it is bigger than its two neighbors in the gradient direction, let `mag_thin[i][j] = mag[i][j]`.
- (17) Plot and return `mag_thin`.

Step 5. Write a function, called `my_linking(mag_thin, orient, tLow, tHigh)`, to perform linking using hysteresis thresholding.

- (18) Create a zero matrix called `result_binary`
- (19) (Forward Scan) Perform a forward scan on rows ($i=0$ to $\text{maxRow}-1$) and columns ($j=0$ to $\text{maxCol}-1$): for each pixel `mag_thin[i][j]` whose value is $\geq t_{\text{High}}$, check if its Right (or BottomRight, or Bottom, or BottomLeft, according to the edge direction) neighboring pixel has a value bigger than t_{Low} . If so, mark that neighboring pixel's value to t_{High} .
- (20) (Backward Scan) Perform a backward scan on rows ($i = \text{maxRow}-1$ to 0) and columns ($j = \text{maxCol}-1$ to 0): for each pixel whose value $\geq t_{\text{High}}$, check if its Left (or TopLeft, or Top, or TopRight) neighboring pixel has a value bigger than t_{Low} , if so, mark that neighboring pixel's value to t_{High} .
- (21) Fill the `result_binary` image: `result_binary[i][j] = 1` if `mag[i][j] $\geq t_{\text{High}}$` .
- (22) Return the `result_binary` image.

Step 6. Compose all the above steps 1-5, to get a function called `my_Canny(img, sigma, tLow, tHigh)`.

Hints:

- a. First test the correctness of your Step 1 ~ 3 using 'TestImg1.jpg' to see whether the gradient orientations are computed correctly.
- b. Adjust t_{Low} and t_{High} to see different detection results. The values should be between 0 and 1.
- c. In this homework, together with your codes, please submit your resultant binary images. Provide a `readme.txt` file, indicating what parameters were you used to generate this resultant images.

Your codes and resultant images are due: 11:59pm, Mar. 3rd.

Submit your homework on Moodle.

If you are late for k ($k < 7$) days, your score will be calculated as:

Final homework score = (The score based on your codes) $\times 0.95^k$

If you are late for more than 7 days, you get no score.